

# Data errors, how to find them?

Edwin de Jonge, Statistics Netherlands

@edwindjonge | [github.com/edwindj](https://github.com/edwindj)



# Who am I?

- ▶ Data scientist / Methodologist at Statistics Netherlands (aka CBS).
- ▶ Author of several R-packages, including `whisker`, `validate`, `errorlocate`, `doceipt`, `daff`, `tableplot`, `ffbase`, `chunked`,  
...
- ▶ Co-author of *Statistical Data Cleaning with applications in R* (2018) (together with @markvdloo)



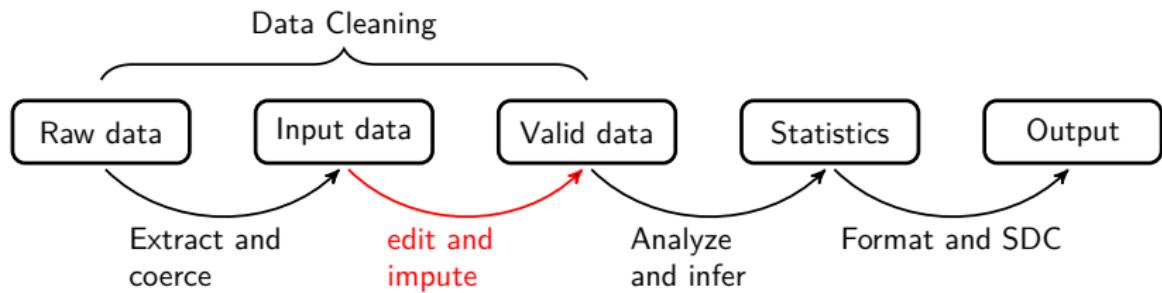
# Data cleaning...

A large part of your job is spent in data-cleaning:

- ▶ getting your data in the right shape (e.g. tidyverse, dplyr)
- ▶ assessing missing data (e.g. VIM, datamaid)
- ▶ checking validity (e.g. validate)
- ▶ locating and removing errors: **errorlocate!**
- ▶ impute values for missing or erroneous data  
(e.g. simputation, VIM, recipes)



# Statistical Value Chain





KEEP  
CALM  
AND  
VALIDATE

# Validation rules?

Package validate allows to:

- ▶ formulate explicit data rule that data must conform to:

```
library(validate)
check_that( data.frame(age=160, driver_license=TRUE),
  age >= 0,
  age < 150,
  if (driver_license == TRUE) age >= 16
)
```



## Explicit validation rules:

- ▶ Give a clear overview what the data must conform to.
- ▶ Can be used to reason about.
- ▶ Can be used to fix/correct data!
- ▶ Find error, and when found correct it.

### Note:

- ▶ Manual fix is error prone, not reproducible and not feasible for large data sets.
- ▶ Large rule set have (very) complex behavior, e.g. entangled rules: adjusting one value may invalidate other rules.



# Error localization

*Error localization is a procedure that points out fields in a data set that can be altered or imputed in such a way that all validation rules can be satisfied.*



## Find the error:

```
library(validate)
check_that( data.frame(age=160, driver_license=TRUE),
  age >= 0,
  age < 150,
  if (driver_license == TRUE) age >= 16
)
```

It is clear that age has an erroneous value, but for more complex rule sets it is less clear.



## Multivariate example:

```
check_that( data.frame( age      = 3
                        , married = TRUE
                        , attends = "kindergarten"
                        )
            , if (married == TRUE) age >= 16
            , if (attends == "kindergarten") age <= 6
            )
```

Ok, clear that this is a faulty record, but what is the error?



## Fellegi Holt formalism:

*Find the minimal (weighted) number of variables that cause the invalidation of the data rules.*

Makes sense! (But there are exceptions...)

Implemented in errorlocate (second generation of editrules).



## Formal description (1)

Rule  $r_i(x)$

A rule a disjunction of atomic clauses:

$$r_i(\mathbf{x}) = \bigvee_j C_i^j(\mathbf{x})$$

with:

$$C_i^j(\mathbf{x}) = \begin{cases} \mathbf{a}^T \mathbf{x} \leq b \\ \mathbf{a}^T \mathbf{x} = b \\ x_j \in F_{ij} \text{ with } F_{ij} \subseteq D_j \\ x_j \notin F_{ij} \text{ with } F_{ij} \subseteq D_j \end{cases}$$

## Rule system:

The rules form a system  $R(\mathbf{x})$ :

$$R_H(\mathbf{x}) = \bigwedge_i r_i$$

If  $R_H(\mathbf{x})$  is true for record  $\mathbf{x}$ , then the record is valid, otherwise one (or more) of the rules is violated.



# Mixed Integer Programming to FH

Each rule set  $R(\mathbf{x})$  can be translated into a mip problem and solved.

$$\begin{aligned} & \text{Minimize } f(\mathbf{x}) = 0; \\ & \text{s.t. } \mathbf{R}\mathbf{x} \leq \mathbf{d} \end{aligned}$$

- ▶  $f(\mathbf{x})$  is the (weighted) number of changed variable:  $\delta_i \in \{0, 1\}$

$$f(\mathbf{x}) = \sum_{i=1}^N w_i \delta_i$$

- ▶  $\mathbf{R}$  contains rules:  $\mathbf{R}_H(\mathbf{x}) \leq \mathbf{d}_H$  and soft constraints:  
 $\mathbf{R}_0(\mathbf{x}, \delta) \leq \mathbf{d}_0$  that try fix the values of  $\mathbf{x}$  to the measured values.

## errorlocate

- ▶ translates your rules automatically into a mip form.
- ▶ Uses lpSolveAPI to solve the problem.
- ▶ contains a small framework for implementing your own error localization algorithms.



## errorlocate::locate\_errors

```
locate_errors( data.frame( age      = 3
                           , married = TRUE
                           , attends = "kindergarten"
                           )
               , validator( if (married == TRUE) age >= 16
                           , if (attends == "kindergarten") age <= 6
                           )
               )$errors

##           age married attends
## [1,] FALSE     TRUE    FALSE
```



## errorlocate::replace\_errors

```
replace_errors(  
    data.frame( age      = 3  
                , married = TRUE  
                , attends = "kindergarten"  
                )  
    , validator( if (married == TRUE) age >= 16  
                , if (attends == "kindergarten") age <= 6  
                )  
)  
  
##   age married      attends  
## 1     3        NA kindergarten
```



## Pipe %>% friendly

The replace\_errors function is pipe friendly:

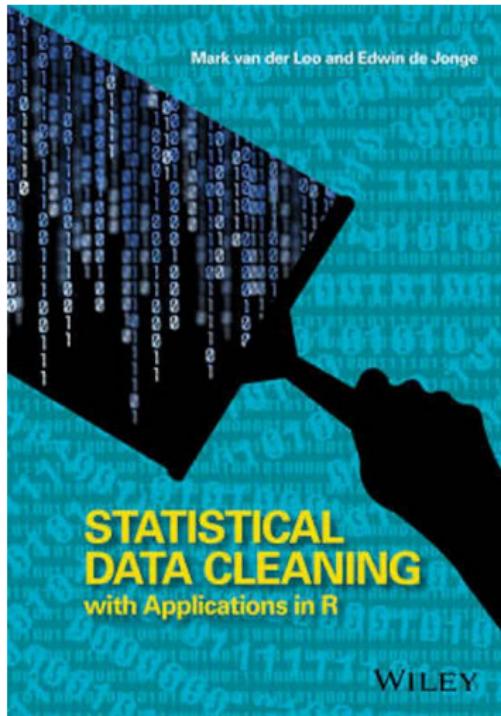
```
rules <- validator(age < 150)

data_noerrors <-
  data.frame(age=160, driver_license = TRUE) %>%
  replace_errors(rules)

errors_removed(data_noerrors) # contains errors removed
```



# Interested?



## SDCR

M. van der Loo and E. de Jonge  
(2018) *Statistical Data Cleaning  
with applications in R* Wiley, Inc.

## errorlocate

- ▶ Available on [CRAN](#)

## More theory?

- ← See book

Thank you for your attention (and enjoy The Hague)!