

Data errors, how to find them?

Edwin de Jonge, Statistics Netherlands

@edwindjonge | github.com/edwindj



Who am I?

- ▶ Data scientist / Methodologist at Statistics Netherlands (aka CBS).
- ▶ Author of several R-packages, including `whisker`, `validate`, `errorlocate`, `doceipt`, `daff`, `tableplot`, `ffbase`, `chunked`,
...
- ▶ Co-author of *Statistical Data Cleaning with applications in R* (2018) (together with @markvdloo)



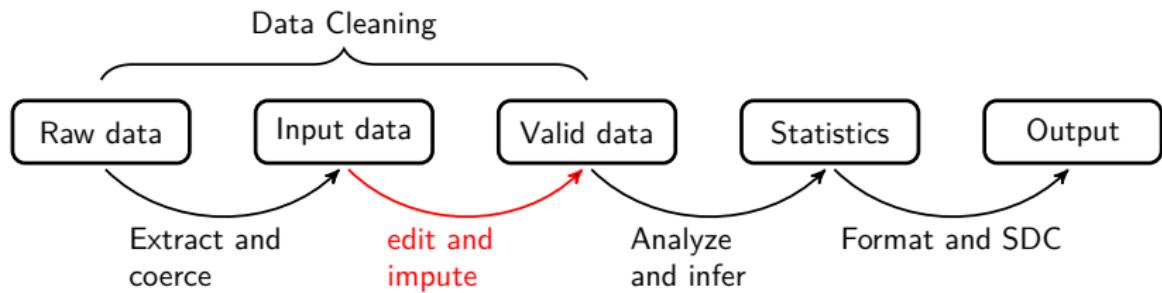
Data cleaning...

A large part of your job is spent in data-cleaning:

- ▶ getting your data in the right shape (e.g. tidyverse, dplyr)
- ▶ assessing missing data (e.g. VIM, datamaid)
- ▶ checking validity (e.g. validate)
- ▶ locating and removing errors: **errorlocate!**
- ▶ impute values for missing or erroneous data
(e.g. simputation, VIM, recipes)



Statistical Value Chain





KEEP
CALM
AND
VALIDATE

Validation rules?

Package validate allows to:

- ▶ formulate explicit data rule that data must conform to:

```
library(validate)
check_that( data.frame(age=160, driver_license=TRUE),
  age >= 0,
  age < 150,
  if (driver_license == TRUE) age >= 16
)
```



Explicit validation rules:

- ▶ Give a clear overview what the data must conform to.
- ▶ Can be used to reason about.
- ▶ Can be used to fix/correct data!
- ▶ Find error, and when found correct it.

Note:

- ▶ Manual fix is error prone, not reproducible and not feasible for large data sets.
- ▶ Large rule set have (very) complex behavior, e.g. entangled rules: adjusting one value may invalidate other rules.



Error localization

Error localization is a procedure that points out fields in a data set that can be altered or imputed in such a way that all validation rules can be satisfied.



Find the error:

```
library(validate)
check_that( data.frame(age=160, driver_license=TRUE),
  age >= 0,
  age < 150,
  if (driver_license == TRUE) age >= 16
)
```

It is clear that age has an erroneous value, but for more complex rule sets it is less clear.



Multivariate example:

```
check_that( data.frame( age      = 3
                        , married = TRUE
                        , attends = "kindergarten"
                        )
            , if (married == TRUE) age >= 16
            , if (attends == "kindergarten") age <= 6
            )
```

Ok, clear that this is a faulty record, but what is the error?



Feligi Holt formalism:

Find the minimal (weighted) number of variables that cause the invalidation of the data rules.

Makes sense! (But there are exceptions...)

Implemented in errorlocate (second generation of editrules).



Formal description (1)

Rule $r_i(x)$

A rule a disjunction of atomic clauses:

$$r_i(\mathbf{x}) = \bigvee_j C_i^j(\mathbf{x})$$

with:

$$C_i^j(\mathbf{x}) = \begin{cases} \mathbf{a}^T \mathbf{x} \leq b \\ \mathbf{a}^T \mathbf{x} = b \\ x_j \in F_{ij} \text{ with } F_{ij} \subseteq D_j \\ x_j \notin F_{ij} \text{ with } F_{ij} \subseteq D_j \end{cases}$$

Rule system:

The rules form a system $R(\mathbf{x})$:

$$R_H(\mathbf{x}) = \bigwedge_i r_i$$

If $R_H(\mathbf{x})$ is true for record \mathbf{x} , then the record is valid, otherwise one (or more) of the rules is violated.



Mixed Integer Programming to FH

Each rule set $R(\mathbf{x})$ can be translated into a mip problem and solved.

$$\begin{aligned} & \text{Minimize } f(\mathbf{x}) = 0; \\ & \text{s.t. } \mathbf{R}\mathbf{x} \leq \mathbf{d} \end{aligned}$$

- ▶ $f(\mathbf{x})$ is the (weighted) number of changed variable: $\delta_i \in \{0, 1\}$

$$f(\mathbf{x}) = \sum_{i=1}^N w_i \delta_i$$

- ▶ \mathbf{R} contains rules: $\mathbf{R}_H(\mathbf{x}) \leq \mathbf{d}_H$ and soft constraints:
 $\mathbf{R}_0(\mathbf{x}, \delta) \leq \mathbf{d}_0$ that try fix the values of \mathbf{x} to the measured values.

errorlocate

- ▶ translates your rules automatically into a mip form.
- ▶ Uses lpSolveAPI to solve the problem.
- ▶ contains a small framework for implementing your own error localization algorithms.



errorlocate::locate_errors

```
locate_errors( data.frame( age      = 3
                           , married = TRUE
                           , attends = "kindergarten"
                           )
               , validator( if (married == TRUE) age >= 16
                           , if (attends == "kindergarten") age <= 6
                           )
               )$errors

##           age married attends
## [1,] FALSE     TRUE    FALSE
```



errorlocate::replace_errors

```
replace_errors(  
    data.frame( age      = 3  
                , married = TRUE  
                , attends = "kindergarten"  
                )  
    , validator( if (married == TRUE) age >= 16  
                , if (attends == "kindergarten") age <= 6  
                )  
)  
  
##   age married      attends  
## 1     3        NA kindergarten
```



Pipe %>% friendly

The replace_errors function is pipe friendly:

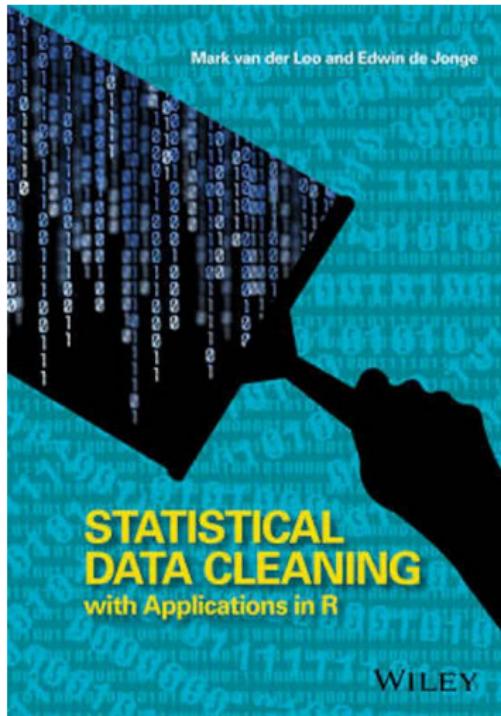
```
rules <- validator(age < 150)

data_noerrors <-
  data.frame(age=160, driver_license = TRUE) %>%
  replace_errors(rules)

errors_removed(data_noerrors) # contains errors removed
```



Interested?



SDCR

M. van der Loo and E. de Jonge
(2018) *Statistical Data Cleaning
with applications in R* Wiley, Inc.

errorlocate

- ▶ Available on [CRAN](#)

More theory?

- ← See book

Thank you for your attention (and enjoy The Hague)!